

CS235102

Data Structures

Chapter 6 Graphs



Chapter 6 Graphs: Outline

- The Graph Abstract Data Type
 - Graph Representations
- Elementary Graph Operations
 - Depth First and Breadth First Search
 - Spanning Tree
- Minimum Cost Spanning Trees
 - Kruskal's, Prim's and Sollin's Algorithm
- Shortest Paths
 - Transitive Closure
- Topological Sorts
 - Activity Networks
 - Critical Paths

The Graph ADT (1/13)

- Introduction

- A graph problem example: Königsberg bridge problem



The Graph ADT (2/13)


- Definitions

- A **graph** G consists of two sets
 - a finite, nonempty set of vertices $V(G)$
 - a finite, possible empty set of edges $E(G)$
- $G(V, E)$ represents a graph
- An **undirected graph** is one in which the pair of vertices in an edge is unordered, $(v_0, v_1) = (v_1, v_0)$
- A **directed graph** is one in which each edge is a directed pair of vertices, $\langle v_0, v_1 \rangle \neq \langle v_1, v_0 \rangle$

tail  **head**



The Graph ADT (3/13)

- Examples for Graph
 - complete undirected graph: $n(n-1)/2$ edges
 - complete directed graph: $n(n-1)$ edges



G_1

complete graph



G_2

incomplete graph



G_3

$$V(G_1) = \{0, 1, 2, 3\}$$

$$V(G_2) = \{0, 1, 2, 3, 4, 5, 6\}$$

$$V(G_3) = \{0, 1, 2\}$$



$$E(G_1) = \{(0, 1), (0, 2), (0, 3), (1, 2), (1, 3), (2, 3)\}$$

$$E(G_2) = \{(0, 1), (0, 2), (1, 3), (1, 4), (2, 5), (2, 6)\}$$

$$E(G_3) = \{<0, 1>, <1, 0>, <1, 2>\}$$

The Graph ADT (4/13)

- Restrictions on graphs
 - A graph may not have an edge from a vertex, i , **back to itself**. Such edges are known as **self loops**
 - A graph may not have multiple occurrences of the same edge. If we remove this restriction, we obtain a data referred to as a **multigraph**



The Graph ADT (5/13)

- Adjacent and Incident
- If (v_0, v_1) is an edge in an undirected graph,
 - v_0 and v_1 are **adjacent**
 - The edge (v_0, v_1) is **incident** on vertices v_0 and v_1





- If $\langle v_0, v_1 \rangle$ is an edge in a directed graph
 - v_0 is **adjacent to** v_1 , and v_1 is **adjacent from** v_0
 - The edge $\langle v_0, v_1 \rangle$ is incident on v_0 and v_1




The Graph ADT (6/13)

- A **subgraph** of G is a graph G' such that $V(G') \subseteq V(G)$ and $E(G') \subseteq E(G)$.



The Graph ADT (6/13)




- A **subgraph** of G is a graph G' such that $V(G') \subseteq V(G)$ and $E(G') \subseteq E(G)$.



The Graph ADT (7/13)


■ Path

- A **path** from vertex v_p to vertex v_q in a graph G , is a **sequence of vertices**, $v_p, v_{i_1}, v_{i_2}, \dots, v_{i_n}, v_q$, such that $(v_p, v_{i_1}), (v_{i_1}, v_{i_2}), \dots, (v_{i_n}, v_q)$ are edges in an undirected graph.
 - A path such as $(0, 2), (2, 1), (1, 3)$ is also written as $0, 2, 1, 3$
- The **length of a path** is the number of edges on it



The Graph ADT (8/13)

- Simple path and cycle
 - **simple path** (simple directed path): a path in which all vertices, except possibly the **first** and the **last**, are **distinct**.
 - A **cycle** is a simple path in which the first and the last vertices are the same.




The Graph ADT (9/13)

- **Connected graph**

- *In an undirected graph G , two vertices, v_0 and v_1 , are connected if there is a path in G from v_0 to v_1*
- An undirected graph is **connected** if, for every pair of distinct vertices v_i, v_j , there is a path from v_i to v_j

The Graph ADT (9/13)

- Connected component
 - A **connected component** of an **undirected graph** is a **maximal connected subgraph**.
 - A **tree** is a graph that is connected and **acyclic** (*i.e., has no cycle*).




The Graph ADT (10/13)

- **Strongly Connected Component**

- A **directed** graph is **strongly connected** if there is a directed path from v_i to v_j and also from v_j to v_i
- A **strongly connected component** is a maximal subgraph that is strongly connected



strongly connected component



The Graph ADT (11/13)


- Degree
 - The **degree** of a vertex is the number of edges incident to that vertex.
- For directed graph
 - **in-degree** (v) : the number of edges that have v as the head
 - **out-degree** (v) : the number of edges that have v as the tail
- If d_i is the degree of a **vertex i** in a graph **G with n vertices and e edges**, the number of edges is

$$e = \left(\sum_{i=0}^{n-1} d_i \right) / 2$$


The Graph ADT (12/13)

- Degree (cont'd)
- We shall refer to a directed graph as a *digraph*. When we use the term *graph*, we assume that it is an *undirected graph*

undirected graph




G_1



G_2

directed graph



G_3


Graph Representations (1/13)

- Adjacency Matrix
- Adjacency Lists
- Adjacency Multilists

Graph Representations (2/13)

- Adjacency Matrix

- Let $G = (V, E)$ be a graph with n vertices.
- The **adjacency matrix** of G is a two-dimensional $n \times n$ array, say adj_mat
- If the **edge** (v_i, v_j) is(not) in $E(G)$, $adj_mat[i][j]=1(0)$



$$\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

G_1




$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

G_3

Graph Representations (2/13)


- Adjacency Matrix

- The adjacency matrix for an undirected graph is symmetric; the adjacency matrix for a digraph need not be symmetric



$$\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

G_1



$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

G_3

Graph Representations (3/13)

- Merits of Adjacency Matrix

- For an undirected graph, the degree of any vertex, i , is its **row sum**:

$$\sum_{j=0}^{n-1} adj_mat[i][j]$$

- For a directed graph, the **row sum** is the **out-degree**, while the **column sum** is the **in-degree**.



$$ind(v_i) = \sum_{j=0}^{n-1} A[j, i] \quad outd(v_i) = \sum_{j=0}^{n-1} A[i, j]$$

Graph Representations (3/13)

- Merits of Adjacency Matrix
 - The complexity of checking edge number or examining if G is connect
 - G is undirected: $O(n^2/2)$
 - G is directed: $O(n^2)$

Graph Representations (5/13)

- Adjacency lists



Graph Representations (4/13)


■ Adjacency lists

- There is one list for each vertex in G . The nodes in list i represent the vertices that are adjacent from vertex i
- For an undirected graph with n vertices and e edges, this representation requires n head nodes and $2e$ list nodes
- C declarations for adjacency lists

```
#define MAX_VERTICES 50 /*maximum number of vertices*/
typedef struct node *node_pointer;
typedef struct node {
    int vertex;
    struct node *link;
};
node_pointer graph[MAX_VERTICES];
int n = 0; /* vertices currently in use */
```

Graph Representations (6/13)

- Sequential Representation of Graph



[0]	9	[12]	3
[1]	11	[13]	0 2
[2]	13	[14]	3
[3]	15	[15]	1 3
[4]	17	[16]	2
[5]	18	[17]	5 4
[6]	20	[18]	4 5
[7]	22	[19]	6
[8]	23	[20]	5 6
[9]	1 0	[21]	7
[10]	2	[22]	6 7
[11]	0 1		

Graph Representations (6/13)

- **Sequential Representation of Graph**
 - Sequentially pack the nodes on the adjacency lists
 - $node[0] \sim node[n-1]$ gives the starting point of the list for vertex i , $0 \leq i < n$
 - $node[n]$ stores “ $n+2e+1$ ”
 - The vertices adjacent from vertex i are stored in $node[node[i]], \dots, node[node[i+1]-1]$, $0 \leq i < n$

Graph Representations (7/13)

- Interesting Operations
 - **degree of a vertex** in an undirected graph
 - # of nodes in adjacency list
 - **# of edges** in a graph
 - determined in $O(n+e)$
 - **out-degree** of a vertex in a directed graph
 - # of nodes in its adjacency list
 - **in-degree** of a vertex in a directed graph
 - traverse the whole data structure

Graph Representations (8/13)

- Finding In-degree of Vertices

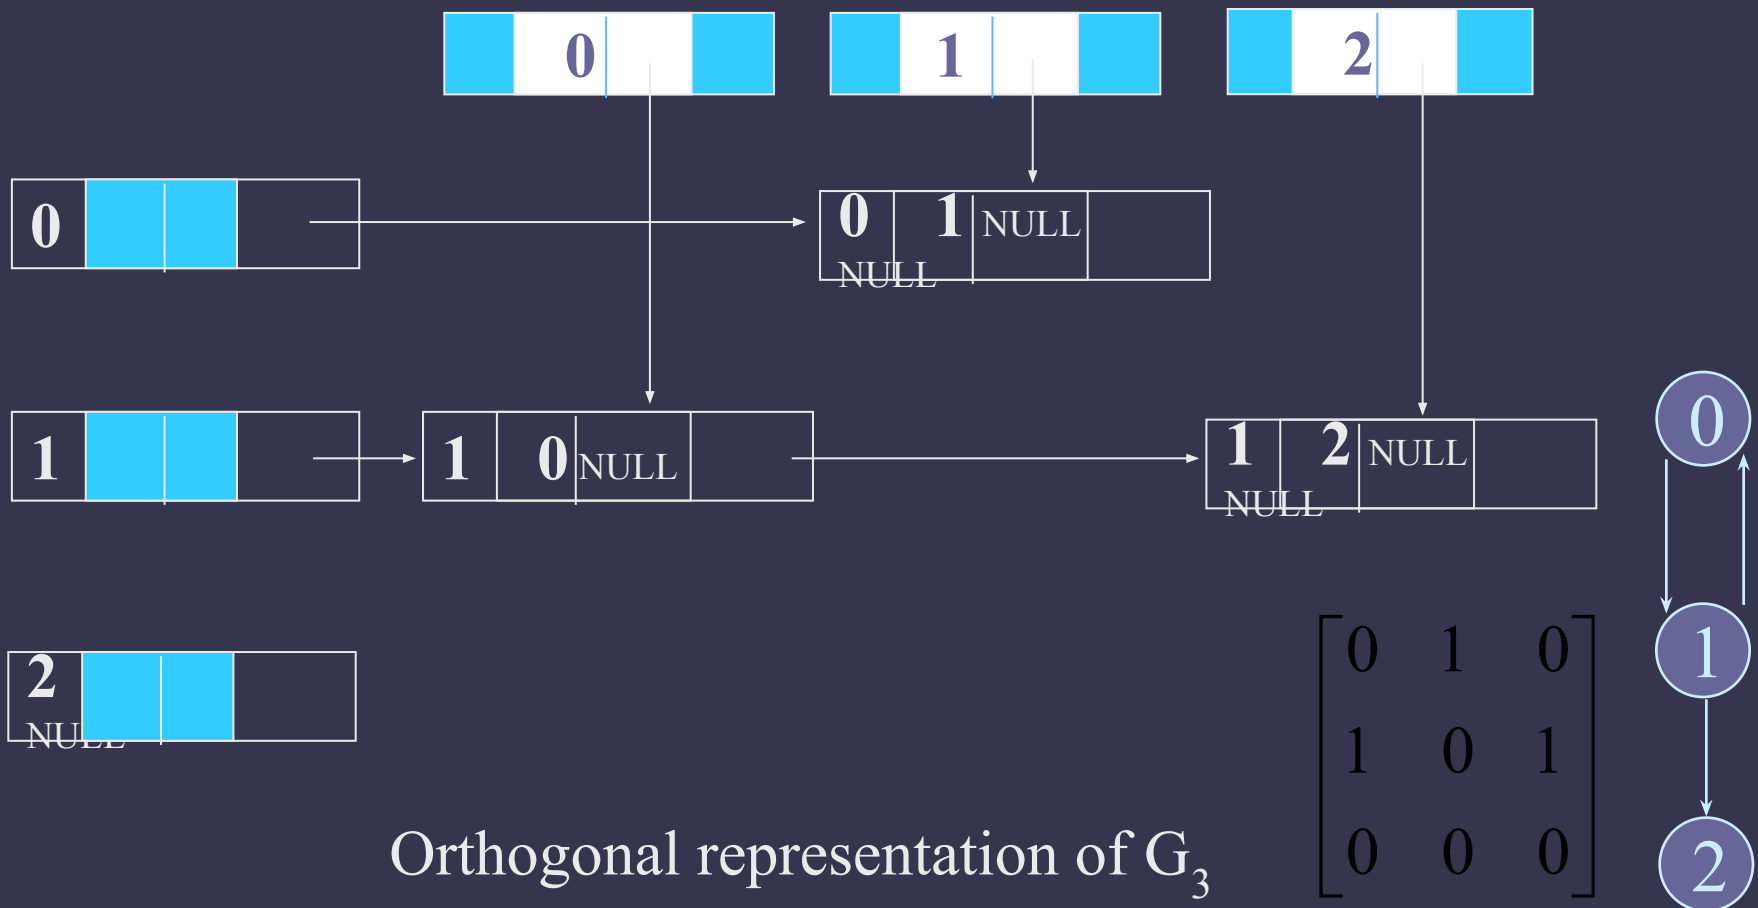


0		1	NULL
1		0	NULL
2		1	NULL

Inverse adjacency list of G_3

Graph Representations (9/13)



- Example of Changing Node Structure



Graph Representations (10/13)

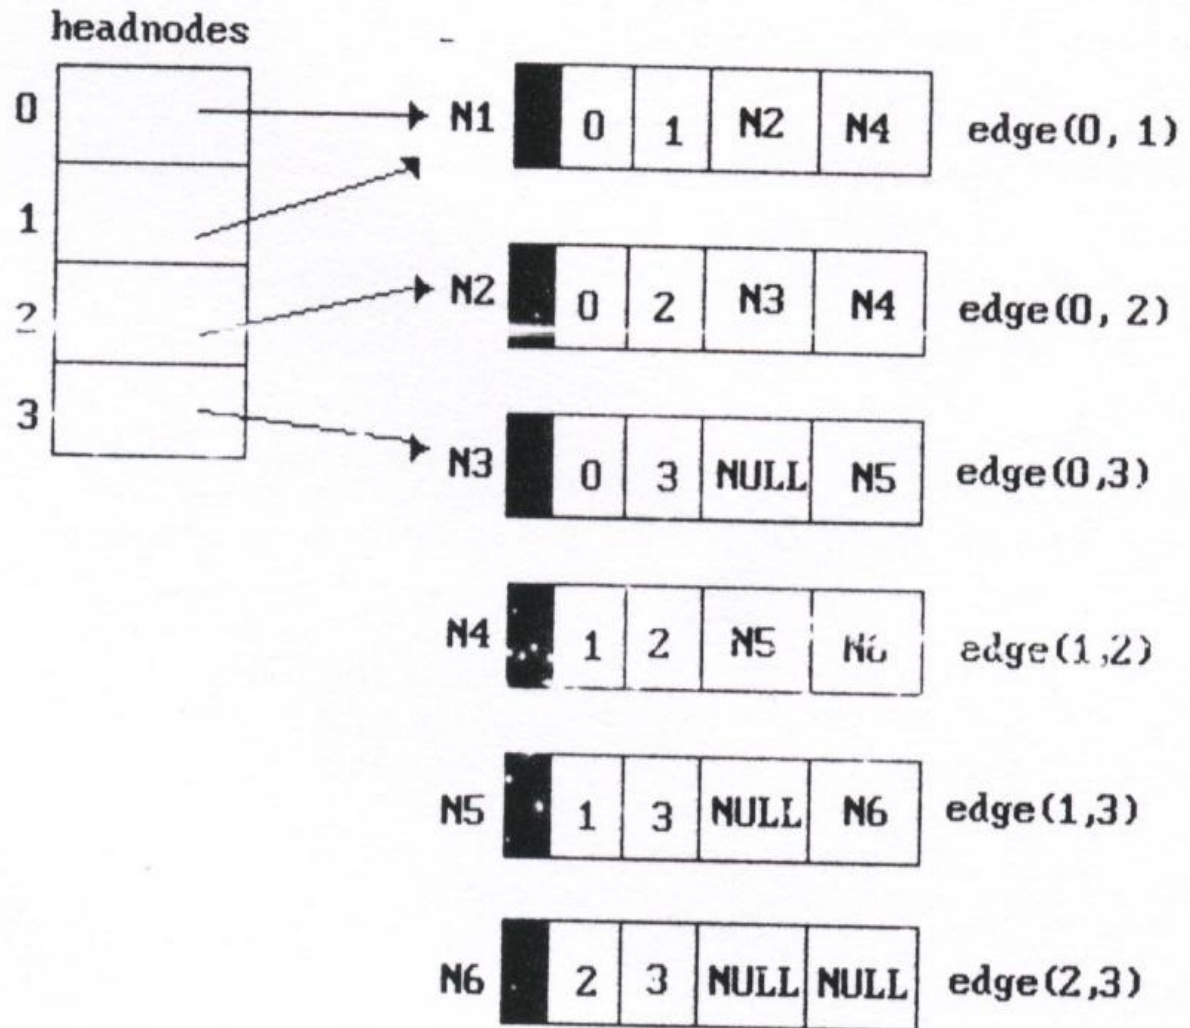

- Vertices in Any Order

Order is of no significance



Graph Representations (12/13)

- Adjacency Multilists



Graph Representations (11/13)

■ Adjacency Multilists

- Lists in which nodes may be shared among several lists. (an edge is shared by two different paths)
- There is exactly one node for each edge.
- This node is on the adjacency list for each of the two vertices it is incident to

marked	vertex1	vertex2	path1	path2
--------	---------	---------	-------	-------

```
typedef struct edge *edge_pointer;
typedef struct edge {
    short int marked;
    int vertex1;
    int vertex2;
    edge_pointer path1;
    edge_pointer path2;
};
edge_pointer graph[MAX_VERTICES];
```

Graph Representations (13/13)

- **Weighted edges**

- The edges of a graph have weights assigned to them.
- These weights may represent as
 - the **distance** from one vertex to another
 - **cost** of going from one vertex to an adjacent vertex.
- **adjacency matrix**: *adj_mat[i][j]* would keep the weights.
- **adjacency lists**: add a *weight* field to the node structure.
- A graph with weighted edges is called a **network**

Graph Operations (1/20)

- Traversal

Given $G=(V,E)$ and vertex v , find all $w \in V$, such that w connects v

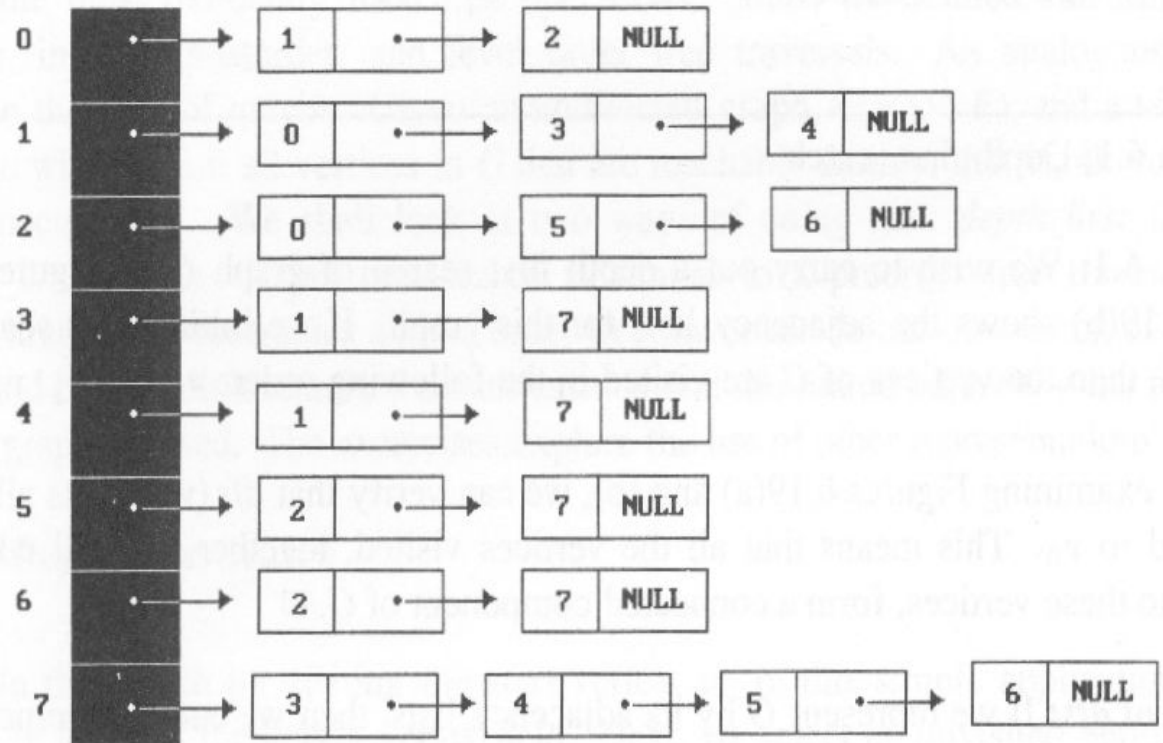

- Depth First Search (DFS): preorder traversal
- Breadth First Search (BFS): level order traversal

- Spanning Trees

- Biconnected Components

Graph Operations (2/20)

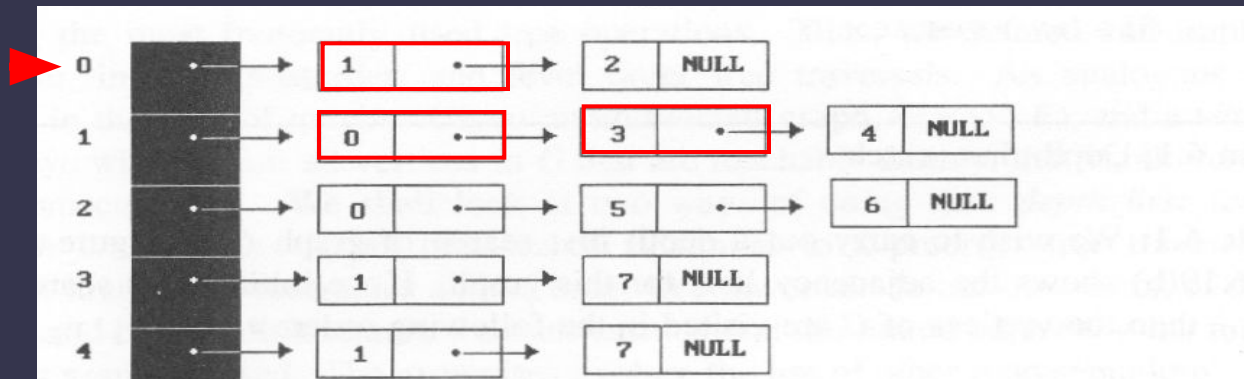
depth first search (DFS): $v_0, v_1, v_3, v_7, v_4, v_5, v_2, v_6$




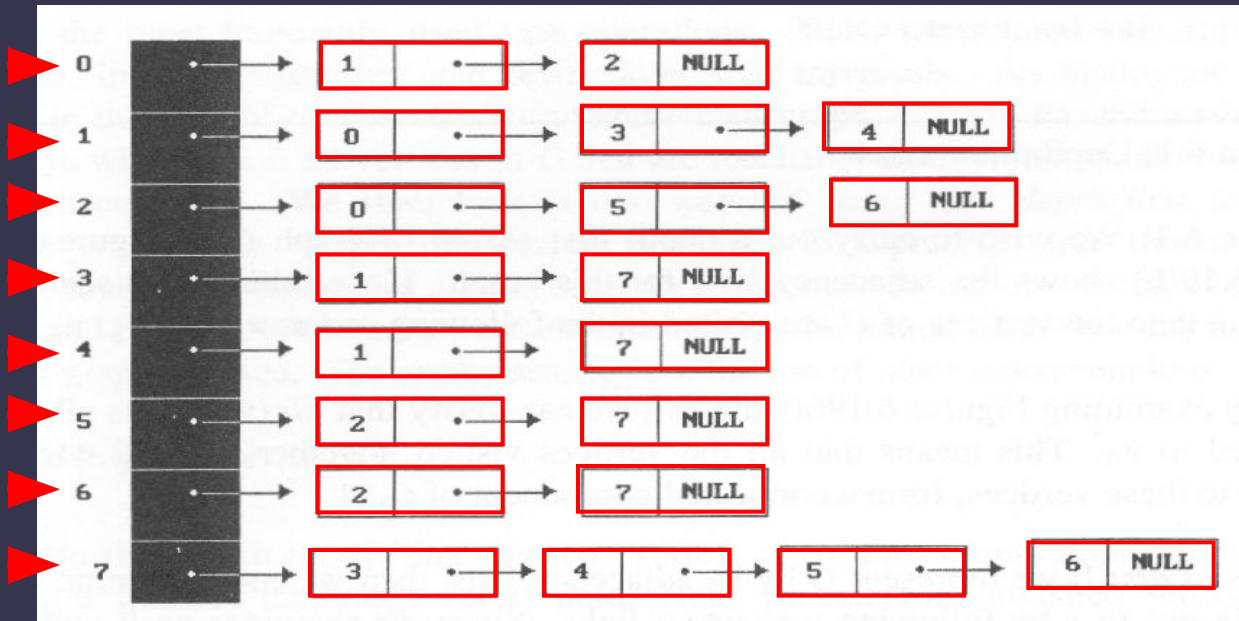
using Adjacency List

■ Depth First Search

```
void dfs(int v)
{
    /* depth first search of a graph beginning with vertex v.*/
    node_pointer w;
    visited[v] = TRUE;
    printf("%5d",v);
    for (w = graph[v]; w; w = w->link)
        if (!visited[w->vertex])
            dfs(w->vertex);
}
```



■ Depth First Search



Data structure
adjacency list: $O(e)$
adjacency matrix: $O(n^2)$

```
void dfs(int v)
{
    /* depth first search of a graph begin
    node_pointer w;
    visited[v] = TRUE;
    printf("%5d",v);
    for (w = graph[v]; w; w = w->link)
        if (!visited[w->vertex])
            dfs(w->vertex);
}
```

```
#define FALSE 0
#define TRUE 1
short int visited[MAX_VERTICES];
```

visited:

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
X	X	X	X	X	X	X	X

output 0 1 3 7 4 5 2 6

:

Graph Operations (4/20)

- Breadth First Search

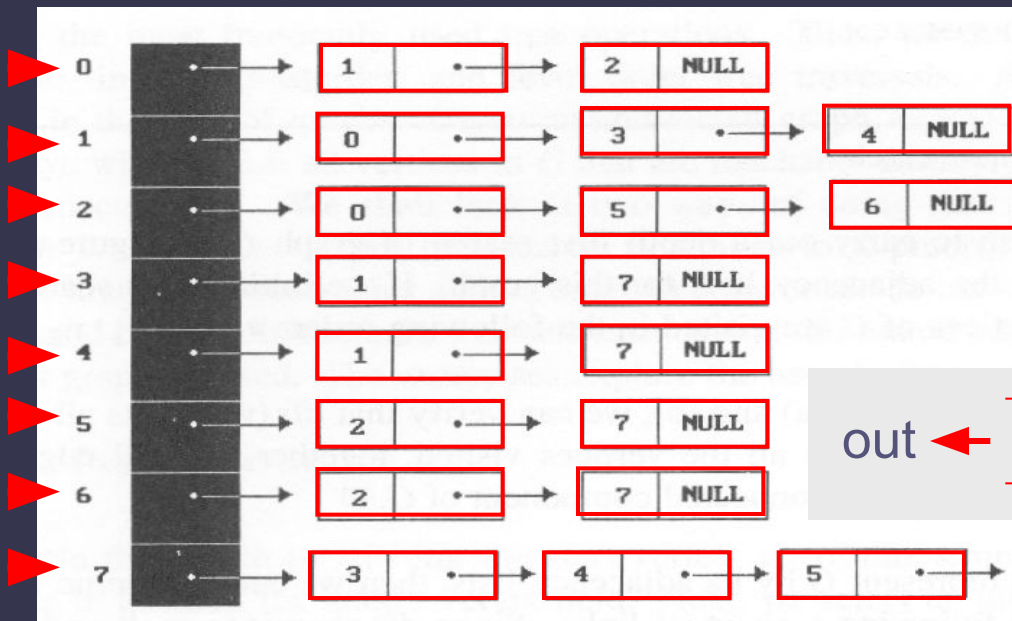
- It needs a queue to implement breadth-first search
- void bfs(int v): breadth first traversal of a graph
 - starting with node v the global array visited is initialized to 0
 - the queue operations are similar to those described in Chapter 4

```
typedef struct queue *queue_pointer;
typedef struct queue {
    int vertex;
    queue_pointer link;
};
void addq(queue_pointer *, queue_pointer *, int);
int deleteq(queue_pointer *);
```

Breadth First Search

```
void bfs(int v)
{
    /* breadth first traversal of a graph, starting with node v
    the global array visited is initialized to 0, the queue
    operations are similar to those described in
    Chapter 4. */
    node_pointer w;
    queue_pointer front, rear;
    front = rear = NULL; /* initialize queue */
    printf("%5d", v);
    visited[v] = TRUE;
    addq(&front, &rear, v);
    while (front) {
        v = deleteq(&front);
        for (w = graph[v]; w; w = w->link)
            if (!visited[w->vertex]) {
                printf("%5d", w->vertex);
                addq(&front, &rear, w->vertex);
                visited[w->vertex] = TRUE;
            }
    }
}
```


Breadth First Search : Example



visited

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
X	X	X	X	X	X	X	X

out ←


0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

 ← in

output 0 1 2 3 4 5 6 7

adjacency list: $O(e)$
adjacency matrix: $O(n^2)$

```
node_pointer w;
queue_pointer front, rear;
front = rear = NULL; /* initialize queue */
printf("%5d", v);
visited[v] = TRUE;
addq(&front, &rear, v);
while (front) {
    v = deleteq(&front);
    for (w = graph[v]; w; w = w->link)
        if (!visited[w->vertex]) {
            printf("%5d", w->vertex);
            addq(&front, &rear, w->vertex);
            visited[w->vertex] = TRUE;
        }
}
```



Graph Operations (6/20)

- Connected components
 - If G is an undirected graph, then one can determine whether or not it is connected:
 - simply making a call to either *dfs* or *bfs*
 - then determining if there is any unvisited vertex

```
void connected(void)
{
    /* determine the connected components of a graph */
    int i;
    for (i = 0; i < n; i++)
        if(!visited[i]) {
            dfs(i);
            printf("\n");
        }
}
```

adjacency list: $O(n+e)$
adjacency matrix: $O(n^2)$



Graph Operations (7/20)

- **Spanning trees**

- **Definition:** A tree T is said to be a *spanning tree* of a connected graph G if T is a subgraph of G and T contains all vertices of G .
- $E(G): T$ (tree edges) + N (nontree edges)
 - T : set of edges used during search
 - N : set of remaining edges




Graph Operations (7/20)

- Spanning trees



Graph Operations (8/20)

- We may use DFS or BFS to create a spanning tree
 - Depth first spanning tree when DFS is used
 - Breadth first spanning tree when BFS is used



Graph Operations (9/20)

- Properties of spanning trees :
 - If a nontree edge (v, w) is introduced into any spanning tree T , then a cycle is formed.
 - A spanning tree is a minimal subgraph, G' , of G such that $V(G') = V(G)$ and G' is connected.
 - We define a minimal subgraph as one with the fewest number of edge
 - A spanning tree has $n-1$ edges

Graph Operations (10/20)


- Assumption: G is an undirected, connected graph
- **Definition:** A vertex v of G is an *articulation point* iff the deletion of v , together with the deletion of all edges incident to v , leaves behind a graph that has at least two connected components.

Graph Operations (10/20)


- **Definition:** A *biconnected graph* is a connected graph that has **no articulation points**.
- **Definition:** A *biconnected component* of a connected graph G is a maximal biconnected subgraph H of G .
 - By maximal, we mean that G contains no other subgraph that is both biconnected and properly contains H .

Graph Operations (11/20)


- Examples of Articulation Points (node 1, 3, 5, 7)



Connected graph




two
connected
components




one connected
graph

Graph Operations (12/20)

- **Biconnected component:**
a maximal biconnected subgraph H
 - no subgraph that is both biconnected and properly contains H



(a) Connected graph




(b) Biconnected components

Graph Operations (13/20)

- Finding the biconnected components
 - By using depth first spanning tree of a connected undirected graph
 - The **depth first number** (*dfn*) outside the vertices in the figures gives the DFS visit sequence
 - If *u* is an ancestor of *v* then $dfn(u) < dfn(v)$

Graph Operations (13/20)

- Finding the biconnected components




result of dfs(3)

Graph Operations (14/20)

- *dfn* and *low*


- Define *low*(*u*): the lowest *dfn* that we can reach from *u* using a path of descendants followed by at most one back edge



- $low(u) = \min\{ dfn(u), \min\{ low(w) \mid w \text{ is a child of } u \}, \min\{ dfn(w) \mid (u, w) \text{ is a back edge} \} \}$

Graph Operations

(14/20)



Graph Operations (15/20)

- Finding an articulation point in a graph:

Any vertex u is an articulation point *iff*

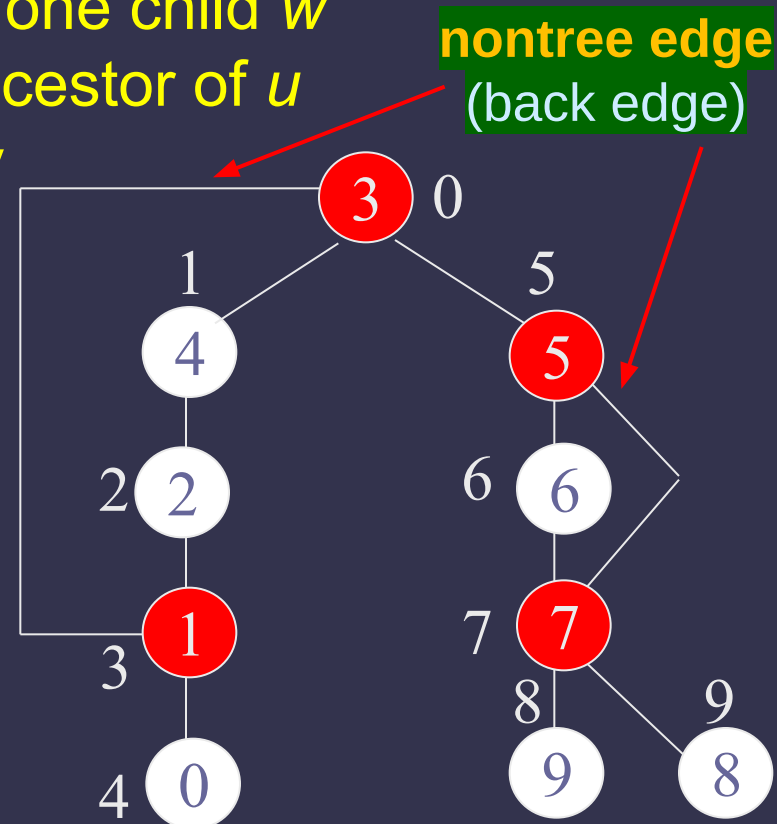
- u is the root of the spanning tree and has two or more children
- u is not the root and has at least one child w such that we cannot reach an ancestor of u using a path that consists of only

(1) w

(2) descendants of w


(3) a single back edge

thus, $low(w) \geq dfn(u)$




Graph Operations (15/20)

- articulation point in a graph:



Graph Operations (15/20)


Vertex	0	1	2	3	4	5	6	7	8	9
<i>dfn</i>	4	3	2	0	1	5	6	7	9	8
<i>low</i>	4	0	0	0	0	5	5	5	9	8



Graph Operations (16/20)

- *dfn* and *low* values for *dfs* spanning tree with root = 3
 - $low(u) = \min\{ dfn(u), \min\{ low(w) \mid w \text{ is a child of } u \}, \min\{ dfn(w) \mid (u, w) \text{ is a back edge} \} \}$

vertex	dfn	low	child	low_child	low:dfn	arti. point
0	4	4 (4,n,n)	null	null	null:4	
1	3	0 (3,4,0)	0	4	$4 \geq 3$	•
2	2	0 (2,0,n)	1	0	$0 < 2$	
3	0	0 (0,0,n)	4,5	0,5	$0,5 \geq 0$	•
4	1	0 (1,0,n)	2	0	$0 < 1$	
5	5	5 (5,5,n)	6	5	$5 \geq 5$	•
6	6	5 (6,5,n)	7	5	$5 < 6$	
7	7	5 (7,8,5)	8,9	9,8	$9,8 \geq 7$	•
8	9	9 (9,n,n)	null	null	null:9	
9	8	8 (8,n,n)	null	null	null:8	



Graph Operations (17/20)

- Determining *dfn* and *low*
 - modify *dfs()* to compute *dfn* and *low* for each vertex of a connected undirected graph


```
void init(void)
{
    int i;
    for (i = 0; i < n; i++) {
        visited[i] = FALSE;
        dfn[i] = low[i] = -1;
    }
    num = 0;
}
```

■ Determining *dfn* and *low* (cont'd)

```
void dfnlow(int u, int v)
{
    node_pointer ptr;
    int w;
    dfn[u] = low[u] = num++;
    for (ptr = graph[u]; ptr; ptr = ptr->link) {
        w = ptr->vertex;
        if (dfn[w] < 0) { /* w is an unvisited vertex */
            dfnlow(w, u);
            low[u] = MIN2(low[u], low[w]);
        }
        else if (w != v)
            low[u] = MIN2(low[u], dfn[w]);
    }
}
```

v is the parent of *u* (if any)

(u, w) is a back edge



Graph Operations (19/20)

- Partition the edges of the connected graph into their biconnected components
 - If $low[w] \geq dfn[u]$, then we have identified a new biconnected component

Graph Operations (19/20)

- We can output all edges in a biconnected component if we **use a stack to save the edges** when we **first encounter** them
- The function *bicon* (Program 6.6) contains the code modified from *dfnlow*, and the same initialization is used

■ Find Biconnected components

```
void bicon(int u, int v)
{
    node_pointer ptr;
    int w,x,y;
    dfn[u] = low[u] = num++;
    for (ptr = graph[u]; ptr; ptr = ptr->link) {
        w = ptr->vertex;
        if (v != w && dfn[w] < dfn[u])
            add(&top,u,w); /* add edge to stack */
        if (dfn[w] < 0) { /* w has not been visited */
            bicon(w,u);
            low[u] = MIN2(low[u],low[w]);
            if (low[w] >= dfn[u]) {
                printf("New biconnected component: ");
                do { /* delete edge from stack */
                    delete(&top, &x, &y);
                    printf(" <%d,%d>",x,y);
                } while (!(x == u) && (y == w));
                printf("\n");
            }
        }
        else if (w != v) low[u] = MIN2(low[u],dfn[w]);
    }
}
```

Find Biconnected components

output <1, 0><1, 3><2, 1><4, 2><3, 4>

: <7, 9>

u <7, 8>

v <7, 5><6, 7>


w <5, 6>

num <3, 5>

num 0

:

:



	0	1	2	3	4	5	6	7	8	9
dfn	-4	-3	-2	-1	-1	-5	-6	-7	-9	-8
low	-4	-3	-2	-1	-1	-5	-6	-7	-9	-8

```
void bicon(int u, int v)
```

```
{
```

```
node_pointer ptr;
```

```
int w,x,y;
```

```
dfn[u] = low[u] = num++;
```

```
for (ptr = graph[u]; ptr; ptr = ptr->link) {
```

```
w = ptr->vertex;
```

back edge or not yet visited

```
if (v != w && dfn[w] < dfn[u])
```

```
add(&top,u,w); /* add edge to stack */
```

```
if (dfn[w] < 0) { /* w has not been visited */
```

```
bicon(w,u);
```

```
low[u] = MIN2(low[u],low[w]);
```

```
if (low[w] >= dfn[u]) {
```

w is a child of u

```
printf("New biconnected component: ");
```

```
do { /* delete edge from stack */
```

```
delete(&top, &x, &y);
```

```
printf(" <%d,%d>",x,y);
```

```
} while (!((x == u) && (y == w)));
```

```
printf("\n");
```

(u, w) is a back edge

```
}
```

```
}
```

```
else if (w != v)
```

```
low[u] = MIN2(low[u],dfn[w]);
```

```
}
```

```
}
```

Minimum Cost Spanning Trees

(1/7)

- Introduction
 - The *cost* of a spanning tree of a weighted, undirected graph is the sum of the costs (weights) of the edges in the spanning tree.
 - A *minimum-cost spanning tree* is a spanning tree of least cost.

Minimum Cost Spanning Trees

(1/7)

- Introduction
 - Three different algorithms can be used to obtain a minimum cost spanning tree.
 - Kruskal's algorithm
 - Prim's algorithm
 - Sollin's algorithm
 - All the three use a design strategy called the greedy method

Minimum Cost Spanning Trees

(2/7)

- **Greedy Strategy**

- Construct an optimal solution in stages
- At each stage, we make the **best decision** (selection) possible at this time.
 - using **least-cost** criterion for constructing minimum-cost spanning trees
- Make sure that the decision will result in a feasible solution
- A feasible solution works within the constraints specified by the problem

Minimum Cost Spanning Trees

(2/7)

- Greedy Strategy
- Our solution must satisfy the following constraints
 - Must use only edges within the graph.
 - Must use exactly $n - 1$ edges.
 - May not use edges that produce a cycle

Minimum Cost Spanning Trees

(3/7)

- Kruskal's Algorithm
 - Build a minimum cost spanning tree T by adding edges to T one at a time
 - **Select the edges** for inclusion in T in **non-decreasing order of the cost**
 - An edge is added to T if it **does not form a cycle**
 - Since G is connected and has $n > 0$ vertices, exactly $n-1$ edges will be selected
 - Time complexity: $O(e \log e)$

Minimum Cost Spanning Trees

(3/7)


- **Theorem 6.1:**

Let G be an undirected connected graph. Kruskal's algorithm generates a minimum cost spanning tree

Minimum Cost Spanning Trees

(4/7)

- Kruskal's Algorithm (cont'd)



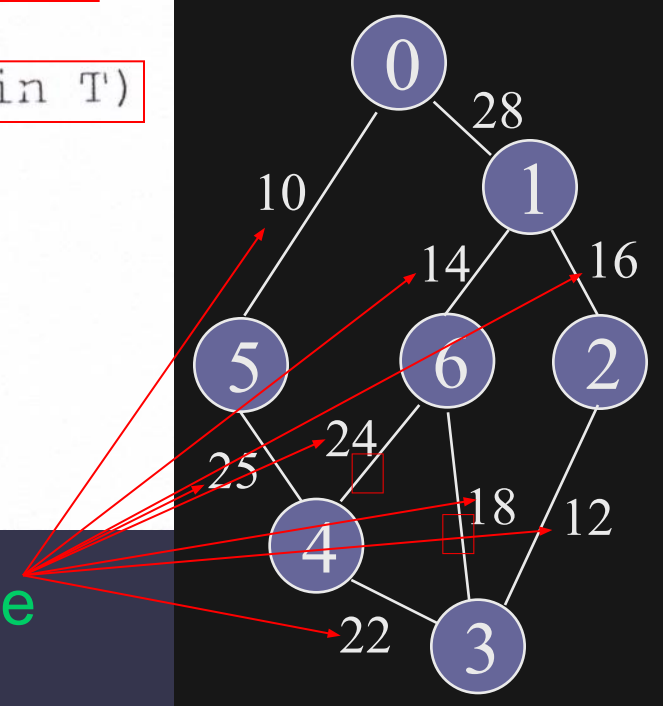
Minimum Cost Spanning Trees

(4/7)

- Kruskal's Algorithm (cont'd)

```
T = {};  
while (T contains less than n-1 edges && E is not empty) {  
    choose a least cost edge (v,w) from E;  
    delete (v,w) from E;  
    if ((v,w) does not create a cycle in T)  
        add (v,w) to T;  
    else  
        discard (v,w);  
}  
if (T contains fewer than n-1 edges)  
    printf("No spanning tree\n");
```

choose



Minimum Cost Spanning Trees

(5/7)


- Prim's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- At each stage, add a least cost edge to T such that the set of selected edges is still a tree.
- Repeat the edge addition step until T contains $n-1$ edges.

Minimum Cost Spanning Trees

(6/7)

- Prim's Algorithm (cont'd)




Minimum Cost Spanning Trees

(6/7)

- Prim's Algorithm (cont'd)


```
T = {};  
TV = {0}; /* start with vertex 0 and no edges */  
while (T contains fewer than n-1 edges) {  
    let (u, v) be a least cost edge such that u ∈ TV and  
    v ∉ TV;  
    if (there is no such edge)  
        break;  
    add v to TV;  
    add (u, v) to T;  
}  
if (T contains fewer than n-1 edges)  
    printf("No spanning tree\n");
```



Minimum Cost Spanning Trees (7/7)

■ Sollin's Algorithm

- Selects several edges for inclusion in T at each stage.
- At the start of a stage, the selected edges forms a spanning forest.
- During a stage we select a minimum cost edge that has exactly one vertex in the tree edge for each tree in the forest.
- Repeat until only one tree at the end of a stage or no edges remain for selection.
 - Stage 1: $(0, 5), (1, 6), (2, 3), (3, 2), (4, 3), (5, 0), (6, 1) \Rightarrow \{(0, 5)\}, \{(1, 6)\}, \{(2, 3), (4, 3)\}$
 - Stage 2: $\{(0, 5), (5, 4)\}, \{(1, 6), (1, 2)\}, \{(2, 3), (4, 3), (1, 2)\}$
 - Result: $\{(0, 5), (5, 4), (1, 6), (1, 2), (2, 3), (4, 3)\}$



Minimum Cost Spanning Trees (7/7)

- Sollin's Algorithm

